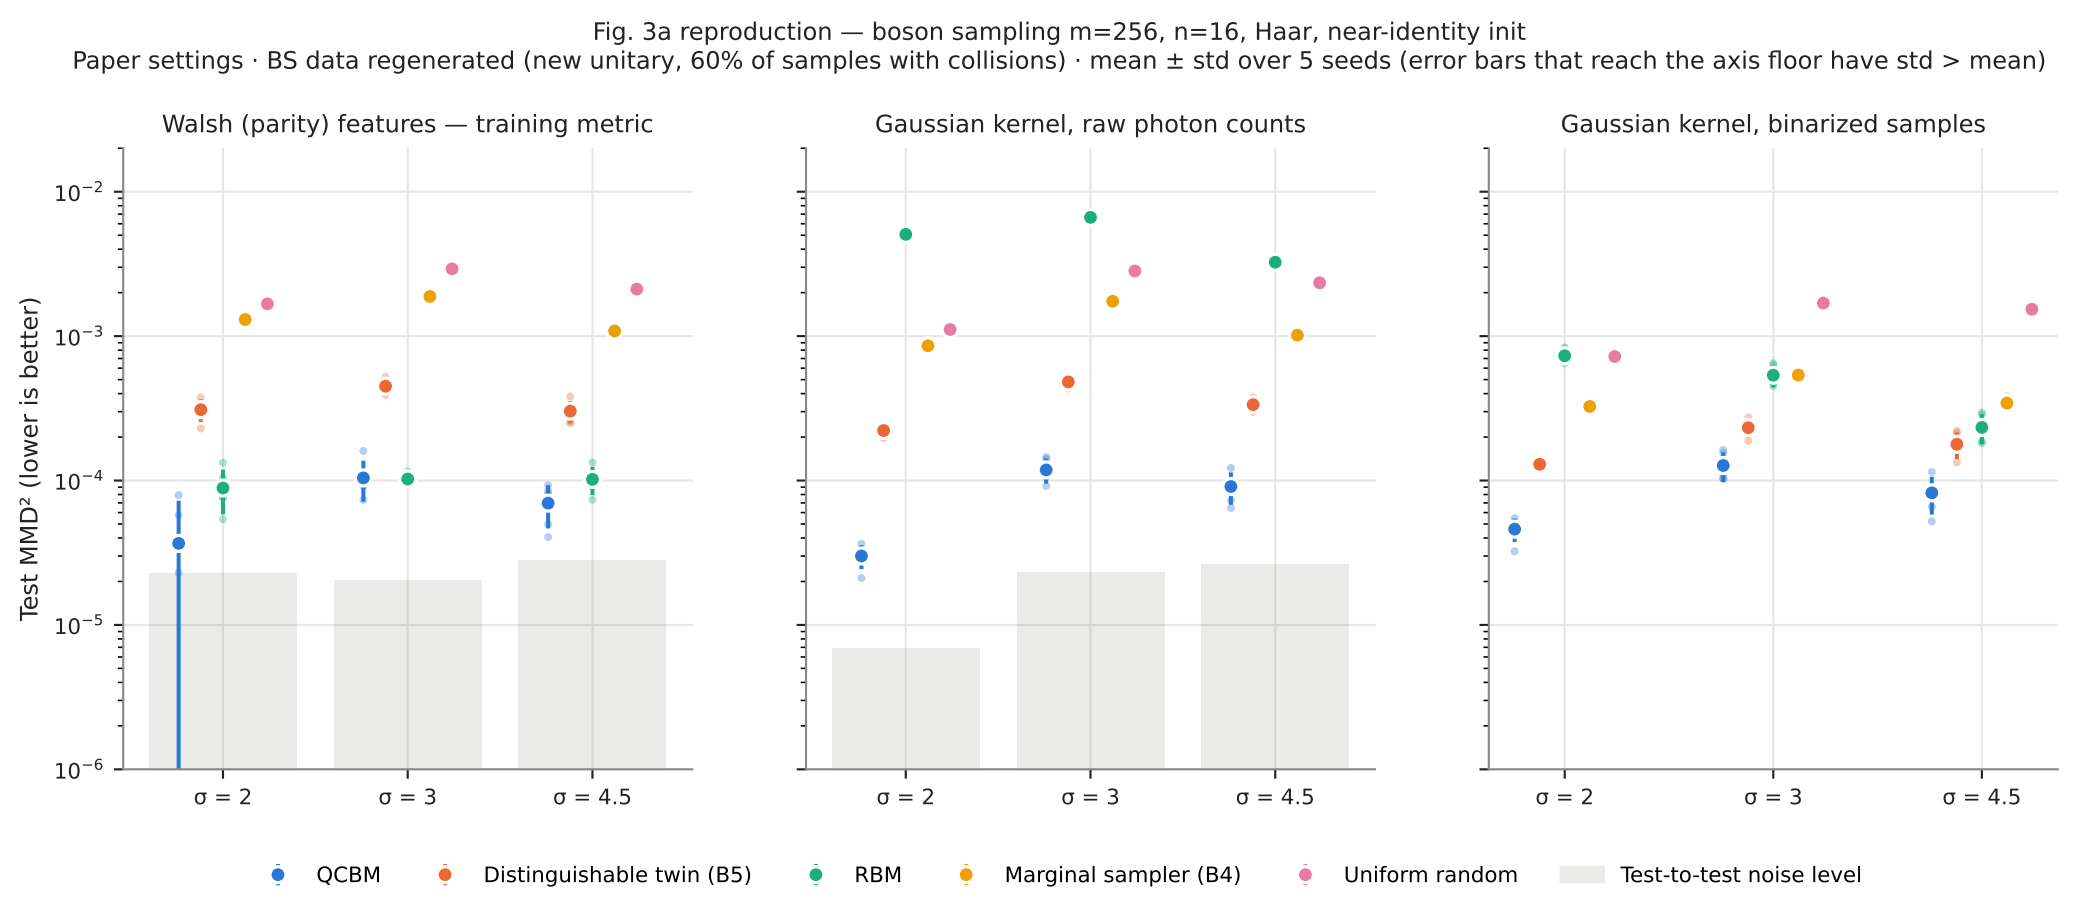
Fig. 3a reproduction — boson sampling m=256, n=16, Haar, near-identity init
Paper settings · BS data regenerated (new unitary, 60% of samples with collisions) · mean ± std over 5 seeds (error bars that reach the axis floor have std > mean)
Walsh (parity) features — training metric
Gaussian kernel, raw photon counts
Gaussian kernel, binarized samples
Test MMD² (lower is better)
10⁻²
10⁻³
10⁻⁴
10⁻⁵
10⁻⁶
σ = 2
σ = 3
σ = 4.5
QCBM
Distinguishable twin (B5)
RBM
Marginal sampler (B4)
Uniform random
Test-to-test noise level