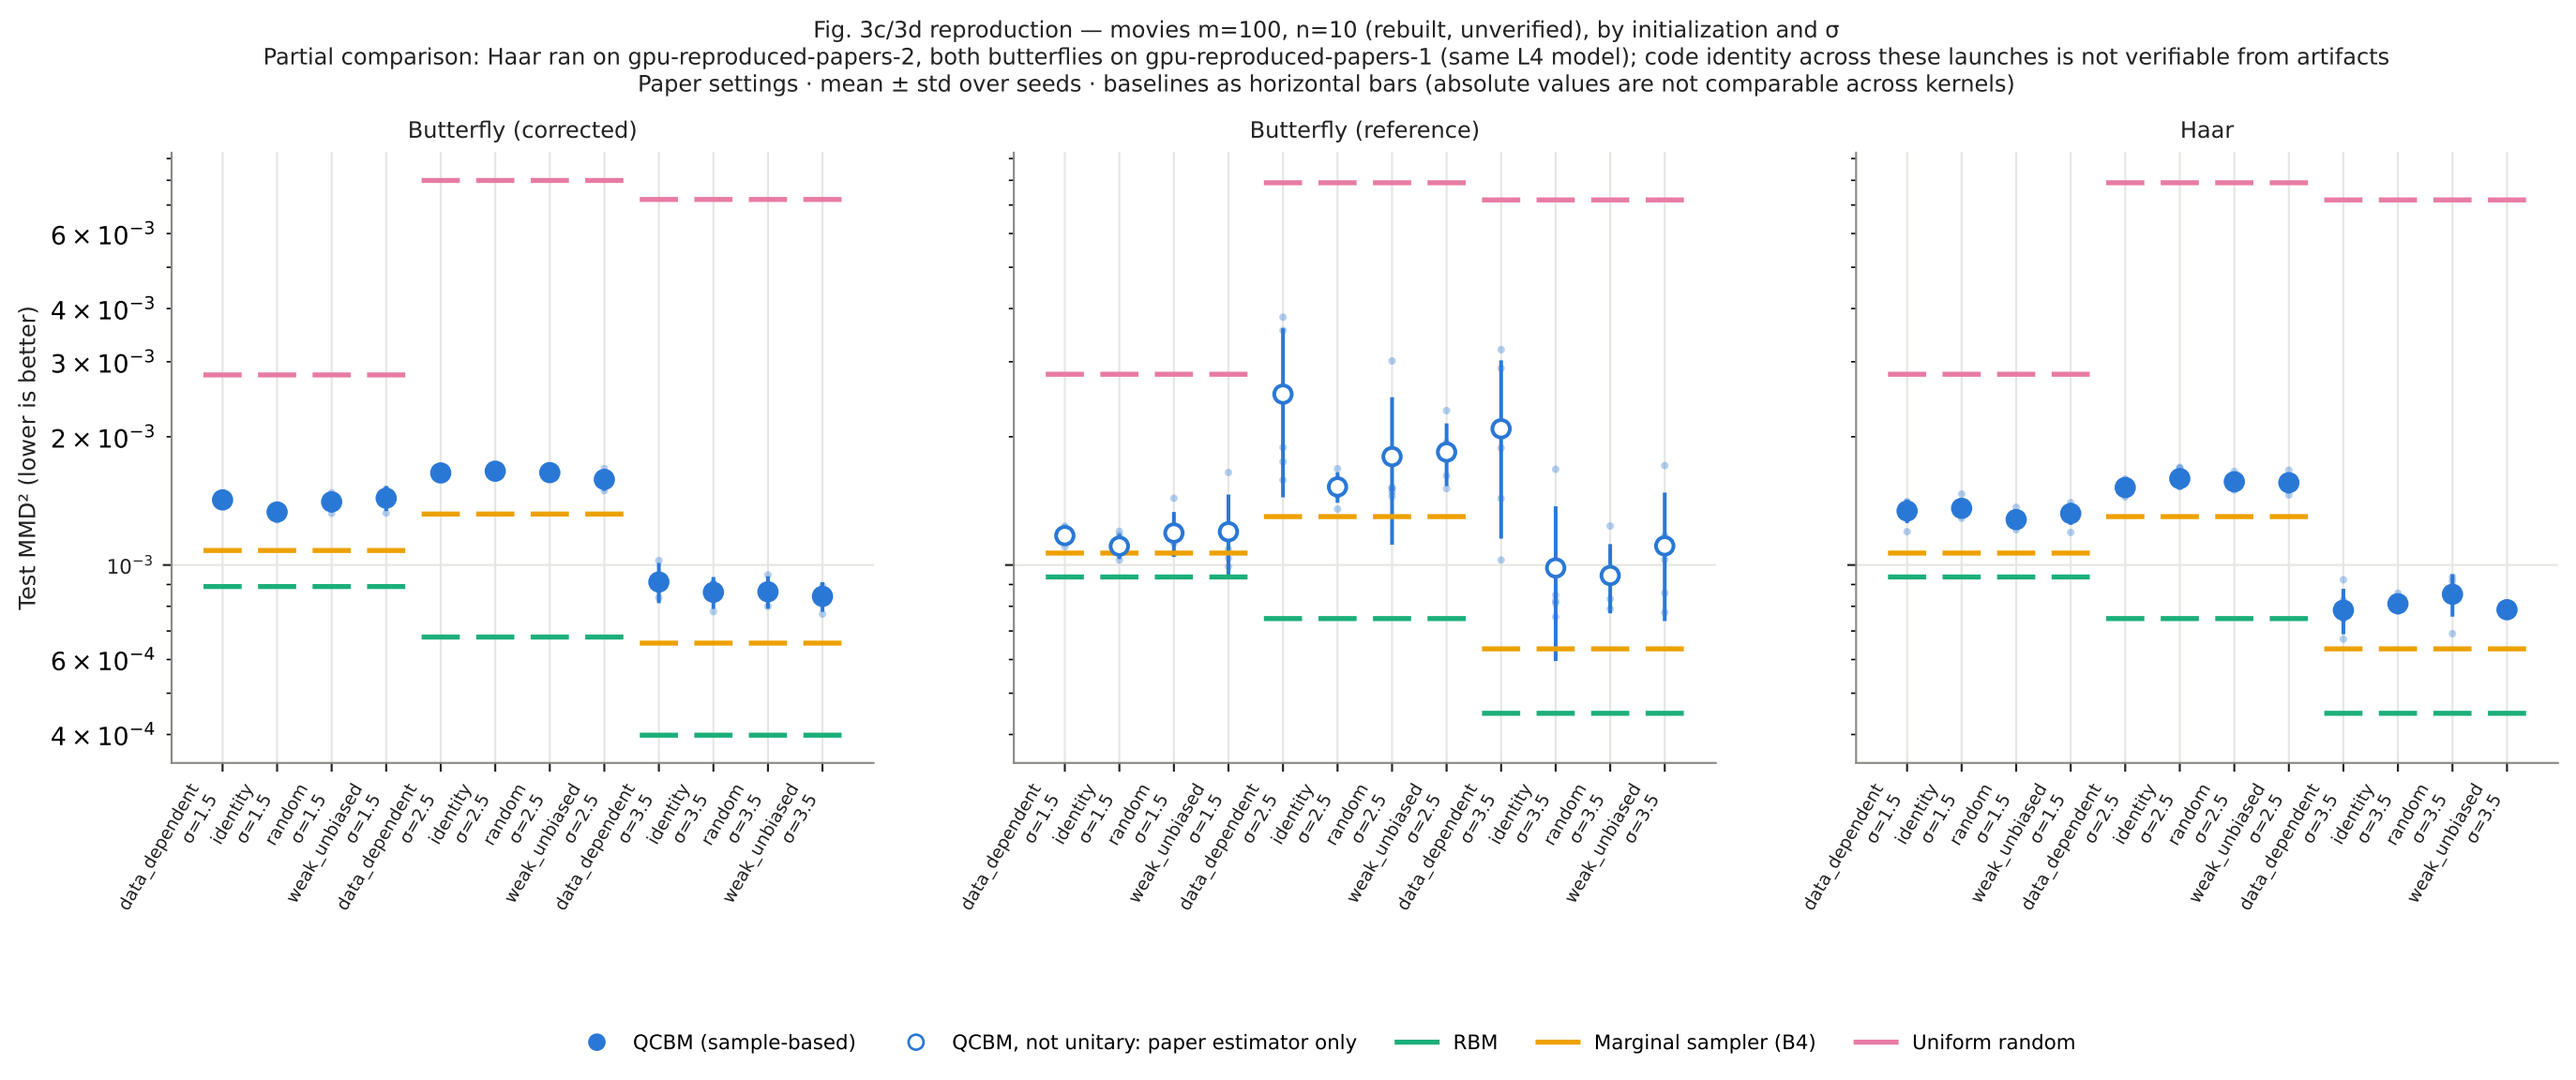
Fig. 3c/3d reproduction — movies m=100, n=10 (rebuilt, unverified), by initialization and σ
Partial comparison: Haar ran on gpu-reproduced-papers-2, both butterflies on gpu-reproduced-papers-1 (same L4 model); code identity across these launches is not verifiable from artifacts
Paper settings · mean ± std over seeds · baselines as horizontal bars (absolute values are not comparable across kernels)
Butterfly (corrected)
Butterfly (reference)
Haar
Test MMD² (lower is better)
6 × 10⁻³
4 × 10⁻³
3 × 10⁻³
2 × 10⁻³
10⁻³
6 × 10⁻⁴
4 × 10⁻⁴
data_dependent σ=1.5
identity σ=1.5
random σ=1.5
weak_unbiased σ=1.5
data_dependent σ=2.5
identity σ=2.5
random σ=2.5
weak_unbiased σ=2.5
data_dependent σ=3.5
identity σ=3.5
random σ=3.5
weak_unbiased σ=3.5
QCBM (sample-based)
QCBM, not unitary: paper estimator only
RBM
Marginal sampler (B4)
Uniform random